

Variables

- CRIM : per capita crime rate by town
- ZN : proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS : proportion of non-retail business acres per town
- CHAS : Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX : nitric oxides concentration (parts per 10 million)
- RM : average number of rooms per dwelling
- AGE : proportion of owner-occupied units built prior to 1940
- AGE : proportion of owner-occupied units built prior to 1940
- DIS : weighted distances to five Boston employment centres
- RAD : index of accessibility to radial highways
- TAX : full-value property-tax rate per \$10,000
- PTRATIO : pupil-teacher ratio by town
- B : $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT : % lower status of the population

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
# from sklearn import datasets
# boston = datasets.load_boston()
# features = pd.DataFrame(boston.data,
# columns=boston.feature_names)
# targets = boston.target
```

```
data_url = "/content/boston.csv"
df = pd.read_csv(data_url)
df.head()
```

| | Unnamed: 0 | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | Price |
|---|------------|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|-------|-------|
| 0 | 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 | 36.2 |

```
del df['Unnamed: 0']
```

```
df.head()
```

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | Price |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|-------|-------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 | 36.2 |

```
##Simple Ensemble Techniques
# Max Voting
# Averaging
```

```
# Weighted Averaging
```

```
# Major Techniques
```

```
# Bagging
```

```
# Boosting
```

```
# Stacking
```

```
df.isnull().sum()
```

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
Price     0
dtype: int64
```

```
#We can also this model
```

```
from sklearn.ensemble import VotingClassifier
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop(['Price'],axis=1)
```

```
y = df.Price
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=32)
```

```
model1 = LogisticRegression(random_state=1)
```


```
model2 = DecisionTreeClassifier(random_state=1)
```

```
model = VotingClassifier(estimators=[('lr', model1), ('dt', model2)], voting='hard')
```

```
# model.fit(X_train,y_train)
```

```
# model.score(X_test,y_test)
```

```
model
```



```
from sklearn.metrics import classification_report
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import RandomForestRegressor
```

```
from sklearn.model_selection import GridSearchCV,RandomizedSearchCV
```

```
X = df.drop(['Price'],axis=1)
```

```
y = df.Price
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=32)
```

```
model = RandomForestRegressor()
```

```
model.fit(X_train,y_train)
```

```
model.score(X_test,y_test)
```

```
0.8739505920595765
```

```
random_grid = {'bootstrap': [True, False],
```

```
                'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None],
```

```
                'max_features': ['auto', 'sqrt'],
```

```
                'min_samples_leaf': [1, 2, 4],
```

```
                'min_samples_split': [2, 5, 10],
```

```
                'n_estimators': [130, 180, 230],
```

```
                'criterion': ['squared_error', 'absolute_error', 'friedman_mse', 'poisson']}]
```

```
# Random search of parameters, using 3 fold cross validation,
```

```
# search across 100 different combinations, and use all available cores
```

```
rf_random = RandomizedSearchCV(RandomForestRegressor(), param_distributions = random_grid,scoring='accuracy',cv=5,verbose=3)
```

```
# Fit the random search model
```

```
rf_random.fit(X_train,y_train)
```

```
rf_random.best_params_
```

```
{'n_estimators': 130,
 'min_samples_split': 10,
 'min_samples_leaf': 2,
 'max_features': 'auto',
 'max_depth': 50,
 'criterion': 'absolute_error',
 'bootstrap': True}
```

```
rf_random.best_estimator_
```

```
RandomForestRegressor
RandomForestRegressor(bootstrap=False, max_depth=40, max_features='sqrt',
 min_samples_split=5, n_estimators=130)
```

```
model = RandomForestRegressor(bootstrap=False, criterion = 'absolute_error', max_depth=40, max_features='sqrt', min_samples_split=5, n_estimators=130)
model.fit(X_train, y_train)
model.score(X_test, y_test)
```

```
0.8882944428271824
```

```
from sklearn.metrics import accuracy_score
y_pred = model.predict(X_test.head(1))
y_pred
```

```
array([16.075])
```